

Multimodal Visual Question Answering with Amazon Berkeley Objects Dataset

Nathan Mathew Verghese
IMT2022022

Divyam Sareen
IMT2022010

Sarvesh Kumar
IMT2022521

I. INTRODUCTION

The task was to build a Visual Question Answering model that takes an image and a related question as input, and responds with a one-word answer. For training, we used the Amazon-Berkeley Objects dataset, along with its metadata, which provided additional details about the items shown in the images.

To approach this, we broke the project down into several key stages: curating the dataset using the Gemini API, setting up a baseline using lightweight yet effective transformer models, and fine-tuning using popular techniques like Low-Rank Adaptation (LoRA). We regularly evaluated the model's performance throughout the training process to track its progress and understand how well it was learning.

II. DATA CURATION:

A. Dataset Overview

In the first stage, we worked with the [1]Amazon-Berkeley Objects Dataset, which includes product images and their associated metadata. We selected the first 20,000 entries from the listings JSON file to create a dataset of 20,000 images, each paired with 3 relevant questions and single-word answers of different difficulty, made by Gemini.

B. Preprocessing

Each entry from the metadata file, `listings.json` file was matched to its corresponding image using a unique image ID, with the mapping provided in the `images.csv` file. For each listing, we first checked whether the resource language was English. If it was, we extracted relevant details such as `item_keywords`, `color`, and `product-type`. These details, along with the associated image, were then passed to the Gemini API to generate meaningful questions.

C. Data Curation Process

For each image and its corresponding metadata extracted from the listing, a prompt (explained later) was constructed and sent to Gemini 1.5 Flash via its API. To handle a large number of requests, multiple API keys were used and rotated throughout the process, along with 25-second delays. Gemini returned a Visual Question-Answer (VQA) pair in JSON format, which was then parsed to verify its validity. Valid responses were saved in a CSV file along with the image ID, image path, and the original listing.

D. The Prompt

To generate Visual Question Answering (VQA) pairs, we used a two-part prompting strategy with Gemini 1.5 Flash: a `system_prompt` to define the task, and a `user_prompt` to pass specific inputs for each image. These prompts were sent to the Gemini API for each image in the dataset.

a) *System Prompt:* The system prompt used is as follows:

You are an expert Visual Question Answering (VQA) dataset creator, specialising in generating questions about product images. Your goal is to create question-answer pairs based **solely** on visual information present in an image.

Your Constraints and Guidelines:

- **Visual Grounding:** All questions must be answerable by looking directly at the image provided. No external knowledge is allowed for answering.
- **Single-Word Answers:** Every answer must be a single, concise word (e.g., colors like "Green", confirmations like "Yes"/"No", materials like "Metal", shapes like "Square").
- **Diverse Questions:** Generate questions covering various visual aspects, including object identification, color, apparent material, shape, key parts, and basic spatial relationships if applicable.
- **Clarity:** Questions should be clear, specific, and unambiguous.
- **Strict Avoidance:** Do NOT generate questions about: subjective qualities, price/brand/origin unless clearly visible as large text, counting numerous small items, reading fine print, or anything requiring external data or numerical output.
- **Output Format:** You MUST output the results as a JSON formatted list of lists. Each inner list must contain exactly two strings: the question and the single-word answer. Example: `[["Question 1", "Answer 1"], ["Question 2", "Answer 2"]]`

b) *User Prompt:* For each image, we dynamically generated the following user prompt using the image ID and extracted keywords:

Image ID: {image_id}
 Metadata Keywords: {keywords}
 Generate 2-3 question-answer pairs based on the image and keywords provided. Follow the JSON output format strictly.

c) *Usage:* These prompts were passed to the Gemini 1.5 Flash model via its API. The model returned a JSON-formatted list of VQA pairs.

III. BASELINE EVALUATION MODELS

We used multiple existing vision-language models to obtain a baseline evaluation on our dataset. The models chosen were Bakllava, BLIP, BLIP-2, and Granite. Each model was selected for its unique architecture and trade-offs between size, performance, and accessibility. For example, Bakllava was chosen because of its high accuracy, but it has 7B model parameters compared to Blip which was giving slightly worse results on our dataset, but had only 400M parameters. To ensure smooth execution and avoid interruptions during long runs on Kaggle, we gracefully handled issues. Additionally, we included a timeout mechanism to halt the notebook safely before Kaggle’s maximum runtime limit of 12 hours was reached. This allowed us to preserve progress and avoid loss of intermediate results.

A. BLIP

Model Description: [2]BLIP (Bootstrapped Language-Image Pretraining) is a base model for Visual Question Answering. It integrates a Vision Transformer (ViT) for extracting features from the image, a BERT-based text encoder for processing the question, and cross-modal attention layers to fuse visual and textual information. BLIP is relatively light-weight and does decently well on some benchmarks. Henceforth, we will be mentioning BLIP as BLIP-VQA-Base.

Implementation: We initialised BLIP with the `Salesforce/blip-vqa-base` weights. GPU acceleration was enabled. For inference, the model processed each image and corresponding question, which was wrapped in a prompt given as `Based on the image, answer the following question with a single word. Question: {question} Answer:.` The entire pipeline was built to load images, tokenise inputs, predict outputs, and clean answers using regular expressions. Results were stored in a CSV file. The script also handled errors like missing or unreadable images and used `tqdm` for tracking progress.

B. BLIP-2

Model Description: BLIP-2 builds upon the original BLIP model by introducing a more modular architecture, including a frozen image encoder and a Q-former (query transformer) for text conditioning. The fused representations are then passed to a large language model (LLM) such as [3]Flan-T5-XL or [4]OPT-2.7B for final answer generation. This decoupled design enables better scalability, flexibility in LLM selection, and improved transferability across tasks.

Implementation: We experimented with two language model backbones for BLIP-2: Flan-T5-XL and OPT-2.7B. The input questions were wrapped the same way as in BLIP-VQA Base.

For the Flan-T5-XL variant, we initialised with the `Salesforce/blip2-flan-t5-xl` weights, which is optimised for accuracy and comprehension across VQA tasks. For the OPT-2.7B variant, we used the `Salesforce/blip2-opt-2.7b` weights, offering a smaller and faster alternative for our needs as Kaggle was posing issues with computation power.

In both cases, the inference loop handled image loading, question formatting, tokenisation, and response cleaning. Results were saved in CSV files after applying regex-based postprocessing.

Code (Flan-T5-XL)

Code (OPT-2.7B)

Quantisation: The OPT-2.7B model was having issues while being loaded onto Kaggle as is. So we used 4-bit quantisation to help us work through this issue. Additionally, we used [5]QLoRA.

C. Bakllava

Model Description: [6]Bakllava is an open-source multimodal model capable of strong performance on vision-language tasks. We used `Bakllava-V1-HF`. It builds on the LLaVA architecture and uses CLIP for vision feature extraction while leveraging a language model backbone. It supports long-context understanding and is optimised for image-grounded tasks such as VQA.

Implementation: The same implementation format was used, similar to the other models. The questions were wrapped in prompts to make the answers one word long. Outputs were parsed from the model’s JSON response and filtered for single-word answers using regex-based cleaning. Processed results were stored in a CSV file for analysis.

Code

D. Granite

Model Description: [7]Granite is IBM’s proprietary vision-language model, designed for tasks like captioning and visual question answering. We used `Granite-Vision-3.1-2B`. It features strong alignment between image and text modalities and is optimised for enterprise-grade applications, balancing performance with inference efficiency.

Implementation: A similar method of implementation was used as in the other models. A wrapper script sent each image and question to the API, parsed the JSON response, validated the answer format, and stored outputs in structured CSV files.

Code

The results and metrics of each model that we ran are given below. Do note that some models only ran half the dataset, some models were quantised because of the memory limitations that we had.

TABLE I: Baseline Model Evaluation Metrics

Model	Dataset	Quantised	EM Accuracy	EM F1	BERT F1
BLIP-VQA-Base	20% of total	No	0.409	0.581	0.882
BLIP-2 (Flan-T5 XL)	100% of total	No	0.488	0.656	0.889
BLIP-2 (OPT-2.7B)	50% of total	Yes (4-bit)	0.508	0.674	0.899
Bakllava	100% of total	No	0.620	0.766	0.920
Granite	50% of total	No	0.673	0.805	0.906

Note: EM stands for Exact Match.

We decided to go ahead with BLIP-VQA-Base, BLIP-2 (Flan-T5 XL), and BLIP-2 (OPT-2.7B) for fine-tuning because these models offered a practical balance between model size and performance. Larger models caused memory issues when loaded on Kaggle, preventing end-to-end training or evaluation. In contrast, these models (BLIP-VQA-Base and BLIP-2 (OPT-2.7B)) were small enough to run reliably within the resource constraints and still produced reasonably strong baseline results, making them suitable candidates for efficient fine-tuning using LoRA. Even BLIP-2 (Flan-T5 XL) which was slightly larger than the other Blip-2 model, was facing memory issues on Kaggle, so we could only train it on limited rows.

IV. FINE TUNING APPROACHES

[8]Low-Rank Adaptation (LoRA) is a parameter-efficient fine-tuning technique that modifies the standard training of transformer-based models by introducing low-rank decomposition matrices. Instead of updating the full weight matrices during backpropagation, LoRA freezes the pre-trained weights and injects trainable matrices $A \in \mathbb{R}^{d \times r}$ and $B \in \mathbb{R}^{r \times k}$ such that the original weight matrix $W_0 \in \mathbb{R}^{d \times k}$ is adapted as:

$$W = W_0 + \Delta W = W_0 + AB$$

This significantly reduces the number of trainable parameters, allowing for efficient fine-tuning even on large models. We used LoRA to adapt our models to our dataset. The configurations are given further.

V. QLoRA

[5]Quantized Low-Rank Adaptation (QLoRA) is an extension of LoRA that further improves memory efficiency by combining parameter-efficient fine-tuning with quantization techniques. While LoRA reduces the number of trainable parameters by injecting low-rank matrices into frozen pre-trained weights, QLoRA goes a step further by operating on quantized (typically 4-bit) versions of these models.

QLoRA utilizes double quantization and paged optimizers to store and train models in reduced precision without significant degradation in performance. This enables the fine-tuning of very large language models on consumer-grade hardware while maintaining competitive accuracy. The key components of QLoRA include:

- **4-bit NormalFloat Quantization:** This technique compresses model weights into 4-bit representations using a

quantisation scheme that preserves essential distributional properties of floating-point numbers.

- **Double Quantization:** Quantization constants themselves are quantized to further reduce memory usage.
- **Paged Optimizers:** Optimizers like Adam are adapted to operate efficiently with quantized parameters, reducing memory fragmentation.

In our work, QLoRA was applied selectively to models that still posed memory challenges after applying LoRA alone. By leveraging quantization, we were able to successfully fine-tune models that would otherwise exceed Kaggle’s memory constraints, such as BLIP-2 (OPT-2.7B) and BLIP-2 (Flan-T5 XL). QLoRA allowed us to retain the advantages of parameter-efficient fine-tuning while fitting larger models into limited memory environments. We only used 4-bit NormalFloat Quantization (NF4).

VI. MODELS, CONFIGURATIONS, AND DATASET SPLITS

Below are the successful configurations that we managed to run:

- **BLIP-VQA-Base:** This model was fine-tuned using LoRA across various rank (r) configurations: 16 (both quantised and non-quantised), 8 (non-quantised), and 32 (non-quantised). An 80-20 train-test split was employed for evaluation.

Code

- **BLIP-VQA-Base (Modified):** To improve performance, we expanded the LoRA target modules beyond the original attention queries and values ([query, value]) to also include the two feed-forward layers and all key, query, value, and output projections: [mlp.fc1, k_proj, v_proj, q_proj, output.dense, mlp.fc2]. This broader coverage allows low-rank adapters to inject task-specific updates into both attention and MLP sublayers, increasing the model’s representational flexibility without substantially raising training cost. As a result, the network better captures multimodal correlations for VQA, yielding higher accuracy and robustness.

Code

- **BLIP-2 (Flan-T5 XL):** Fine-tuned using LoRA with $r = 8, 16$ (tried) and quantisation enabled. It was trained and tested on a limited (300) rows due to computational bounds.

Code

- **BLIP-2 (OPT-2.7B)**: Fine-tuned using LoRA with $r = 8, 16$ (tried) and quantisation enabled. The dataset split used was 50-50.

Code

The top models on the test split are given in the table below. The best-performing model is highlighted in bold.

TABLE II: Evaluation Metrics for Fine-Tuned Models Using LoRA

Model	Train-Test Split	Quantised (NF4)	LoRA r	EM Acc.	EM F1	BERT F1
BLIP-VQA-Base	80-20	No	8	0.485	0.653	0.902
BLIP-VQA-Base (M)	80-20	No	16	0.570	0.726	0.921
BLIP-VQA-Base	80-20	No	16	0.528	0.691	0.899
BLIP-VQA-Base (M)	80-20	Yes	16	0.562	0.720	0.922
BLIP-VQA-Base	80-20	No	32	0.524	0.688	0.901
BLIP-2 (Flan-T5 XL)	300 questions tested	Yes	8	0.367	0.537	0.775
BLIP-2 (OPT-2.7B)	50-50	Yes	16	0.508	0.673	0.899

Note: EM stands for Exact Match. (M) signifies that the target modules were modified.

VII. EVALUATION METRICS

We used 3 main metrics (but did not limit to just this).

- **Exact-Match Accuracy**: This checks if the predicted words exactly match the ground-truth words after removing special symbols and ignoring capitalisation. Since we predict only one word, the accuracy will be equal to the recall.
- **Exact-Match F1 score**: Here, we assume precision is 1, meaning everything the model predicted is correct. The recall is then the fraction of the ground-truth tokens that appear in the prediction:

$$\text{Recall} = \frac{\text{Number of correctly predicted tokens}}{\text{Total tokens in the ground truth}}$$

Since precision $P = 1$, the F1 score simplifies to:

$$F1 = 2 \times \frac{P \times R}{P + R} = \frac{2R}{1 + R}$$

- **BERT F1 score**: Unlike exact matching, BERT F1 uses BERT’s contextual embeddings to compare predicted and true tokens. Instead of only checking if tokens are exactly the same, it measures how similar they are in meaning.

This makes BERT F1 better for cases where the prediction is close in meaning but not a perfect word-for-word match.

- **Precision** measures how many predicted tokens have a good semantic match in the ground truth.
- **Recall** measures how much of the ground truth is captured by the prediction’s meaning.
- **F1** combines precision and recall into a single score reflecting overall semantic overlap.

Because it looks at meaning rather than exact words, BERT F1 usually gives a more useful score when evaluating natural language tasks.

In addition to Bert Score, to evaluate the quality of the generated answers, we used multiple other metrics: ROUGE, Levenshtein Normalized Similarity, and Sentence-BERT Cosine Similarity.

- **ROUGE1**: Refers to the overlap of unigrams (each word) between the system and reference summaries. A higher ROUGE1 score means the prediction is more similar in wording to the reference.
- **Levenshtein Normalized Similarity**: Calculates how many character-level edits (insertions, deletions, or substitutions) are needed to change the predicted sentence into the reference. A higher score indicates the two strings are more similar.
- **Sentence-BERT Cosine Similarity**: Measures the similarity in meaning between the predicted and reference answers using sentence embeddings. A higher cosine similarity shows that both answers are close in meaning.

Using a mix of these metrics helps us evaluate both the surface-level and deeper semantic similarity between the generated and expected answers.

VIII. CONCLUSION

In this project, we developed a robust multimodal Visual Question Answering pipeline using the Amazon-Berkeley Objects dataset. We curated a high-quality dataset leveraging Gemini API-generated question-answer pairs grounded in image metadata and established strong baselines with various vision-language models, including BLIP and BLIP-2. Our approach highlights the power of combining structured data, automated prompt engineering, and model fine-tuning techniques like LoRA to address real-world VQA challenges effectively. Future work could explore improved reasoning models and dataset expansion for even richer understanding,

and if we had stronger compute, we could have tried larger models as well.

Among all the fine-tuned models evaluated, the **BLIP-VQA-Base (M)** model with a train-test split of 80-20, no quantisation, and a LoRA rank of 16 achieved the best overall results. It obtained the highest scores across all key metrics: an Exact Match Accuracy (EM Acc.) of 0.570, EM F1 of 0.726, and BERT F1 of 0.921. These results suggest that modifying the target modules and choosing an appropriate LoRA configuration (in this case, $r = 16$) can significantly improve the model's ability to generate accurate and semantically meaningful answers.

Link to github: <https://github.com/nathanmathewv/Multimodal-VQA>

Link to dataset: <https://www.kaggle.com/datasets/nathanmathew/images-with-vqas>

IX. REFERENCES

REFERENCES

- [1] Amazon Berkeley Objects (ABO). Available at: <https://amazon-berkeley-objects.s3.amazonaws.com/index.html>
- [2] BLIP GitHub repository. Available at: <https://github.com/salesforce/BLIP>.
git
- [3] BLIP2 Flan-T5-XL model on Hugging Face. Available at: <https://huggingface.co/Salesforce/blip2-flan-t5-xl>
- [4] BLIP2 OPT-2.7B model on Hugging Face. Available at: <https://huggingface.co/Salesforce/blip2-flan-t5-xl>
- [5] The first paper on QLoRA is given here: <https://arxiv.org/abs/2305.14314>
- [6] Bakllava model on Hugging Face. Available at: <https://huggingface.co/llava-hf/bakLlava-v1-hf>
- [7] Granite model on Hugging Face. Available at: <https://huggingface.co/ibm-granite/granite-vision-3.2-2b>
- [8] Check LoRA out here: <https://huggingface.co/docs/diffusers/main/en/training/lora>